

# Deployment strategy in kubernetes

1. [Pods](#)
2. [Replication Controller](#)
3. [ReplicaSet](#)
4. [Deployment](#)
5. [DaemonSet](#)

## 1. **Pods** – Let's perform some tasks using Pods strategy.

Alright! Now, we will create a **Pod** using the **file method** in Kubernetes instead of using a Deployment. This approach is useful when you want full control over individual Pods without using a ReplicaSet or Deployment.

### Steps to Create a Pod Using a YAML File

1. **Create a YAML file** defining the Pod.
2. **Apply the YAML file** using kubectl.
3. **Verify the Pod status.**

#### 1. **Create a YAML file** defining the Pod.

```
vim pod.yml
```

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
    - name: test
      image: nginx
```

```
kubectl create -f pod.yml
```

```
root@master-node:~# vim pod.yml
root@master-node:~# kubectl create -f pod.yml
pod/mypod created
root@master-node:~# |
```

```
root@master-node:~# kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
mypod    1/1     Running   0          56s
root@master-node:~# |
```

```
kubectl describe pod mypod
```

```
Events:
Type  Reason  Age   From           Message
----  -----  --   --            --
Normal Scheduled  100s  default-scheduler  Successfully assigned default/mypod to worker1-node.
Normal Pulling   99s  kubelet        Pulling image "nginx"
Normal Pulled    99s  kubelet        Successfully pulled image "nginx" in 259ms (259ms including waiting). Image size: 72188133 bytes.
Normal Created   98s  kubelet        Created container test
Normal Started   98s  kubelet        Started container test
root@master-node:~# |
```

```
kubectl delete pod mypod --force
```

```
root@master-node:~# kubectl delete pod mypod --force
Warning: Immediate deletion does not wait for confirmation that the running resource has been terminated. The resource may continue to run on the cluster indefinitely.
pod "mypod" force deleted
root@master-node:~# |
```

```
root@master:~# kubectl delete -f pod.yaml
pod "rupesh" deleted
root@master:~#
root@master:~# kubectl get pods
```

**Next – How we can run a command inside the container**

```
kubectl exec mypod -- hostname
```

```
root@master-node:~# kubectl exec mypod -- hostname
mypad
root@master-node:~# kubectl exec mypod -- pwd
/
root@master-node:~# kubectl exec mypod -- ls
bin
boot
dev
docker-entrypoint.d
docker-entrypoint.sh
etc
```

### Login into the pod.

```
kubectl exec mypod -it -- bash
```

```
root@master-node:~#
root@master-node:~# kubectl exec mypod -it -- bash
root@mypad:/# pwd
/
root@mypad:/# |
```

```
root@master-node:~# kubectl logs mypod
/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Sourcing /docker-entrypoint.d/15-local-resolvers.envsh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/02/19 13:11:13 [notice] 1#1: using the "epoll" event method
2025/02/19 13:11:13 [notice] 1#1: nginx/1.27.4
2025/02/19 13:11:13 [notice] 1#1: built by gcc 12.2.0 (Debian 12.2.0-14)
2025/02/19 13:11:13 [notice] 1#1: OS: Linux 6.8.0-1021-aws
2025/02/19 13:11:13 [notice] 1#1: getrlimit(RLIMIT_NOFILE): 1048576:1048576
2025/02/19 13:11:13 [notice] 1#1: start worker processes
2025/02/19 13:11:13 [notice] 1#1: start worker process 29
root@master-node:~# |
```

### Copy content inside the pod

```
kubectl cp abc mypod:/
```

```
root@master-node:~# kubectl cp abc mypod:/  
root@master-node:~#  
root@master-node:~# kubectl exec mypod -- ls  
abc  
bin  
boot
```

## How we can assign labels to pods

```
vim pod.yml
```

```
apiVersion: v1  
kind: Pod  
metadata:  
    name: mypod  
    labels:  
        app: wordpress  
        env: prod  
spec:  
    containers:  
        - name: nam  
          image: nginx  
          ports:  
            - containerPort: 80
```

```
kubectl create -f pod.yml
```

```
kubectl describe pod mypod
```

```
root@master-node:~# kubectl create -f pod.yml
pod/mypod created
root@master-node:~# kubectl describe pod mypod
Name:           mypod
Namespace:      default
Priority:       0
Service Account: default
Node:           worker1-node.com/172.31.17.146
Start Time:     Wed, 19 Feb 2025 14:00:20 +0000
Labels:         app=wordpress
                env=prod
Annotations:   <none>
Status:        Running
IP:            10.46.0.1
IPs:
    IP: 10.46.0.1
```

```
kubectl get pods --show-labels
```

```
root@master-node:~# kubectl get pods --show-labels
NAME    READY   STATUS    RESTARTS   AGE    LABELS
mypod  1/1     Running   0          4m45s   app=wordpress,env=prod
root@master-node:~# |
```

## Why Do We Assign Labels to Pods in Kubernetes?

Labels in Kubernetes are **key-value pairs** assigned to objects (like Pods) to organize, identify, and manage them efficiently. They help in **grouping, selecting, and filtering** resources.

### Key Reasons for Using Labels in Pods

#### 1. Logical Grouping of Pods

- Helps categorize Pods based on environment, application, or role.
- Example:

```
labels:
  app: webserver
  environment: production
```

- o Now, all Pods running the **webserver** in the **production** environment can be easily managed together.

## 2. Service Discovery & Load Balancing

- Services use labels to select the right Pods for traffic routing.
- Example:

```
selector:  
  app: webserver
```

- o This ensures that the **Service** only routes traffic to Pods with the app: webserver label.

## 3. Managing Deployments & Scaling

- Labels help **ReplicationControllers**, **ReplicaSets**, and **Deployments** manage specific sets of Pods.
- Example:

```
selector:  
  matchLabels:  
    app: database
```

- o A **ReplicaSet** will ensure the required number of database Pods always runs.

## 4. Easy Filtering & Querying

- Helps in searching for specific Pods quickly.
- Example:

```
kubectl get pods -l app=webserver
```

- o Lists only the Pods labeled as webserver.

## 5. Rollouts & Canary Deployments

- Helps in **progressive deployment** and testing new versions before full rollout.
- Example:

```
labels:  
  app: webserver  
  version: v1
```

- o Later, a v2 version can be tested by selecting only version: v2 Pods.

## 2. ReplicationController – Let's perform some tasks using replication controller

```
# vim rc.yaml  
  
apiVersion: v1  
  
kind: ReplicationController  
  
metadata:  
  name: myapp  
  
spec:  
  replicas: 10  
  
  template:  
    metadata:  
      name: xyz  
  
    labels:  
      app: webserver
```

```
spec:
```

```
  containers:
```

```
    - name: sdfdf
```

```
      image: httpd
```

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: myapp
spec:
  replicas: 10
  template:
    metadata:
      name: xyz
      labels:
        app: webserver
    spec:
      containers:
        - name: sdfdf
          image: httpd
```

```
kubectl create -f rc.yaml
```

```
kubectl get rc
```

```
kubectl get pods
```

```
root@master-node:~# kubectl create -f rc.yaml
replicationcontroller/myapp created
root@master-node:~#
root@master-node:~# kubectl get rc
NAME    DESIRED   CURRENT   READY   AGE
myapp   10        10        10      19s
root@master-node:~#
root@master-node:~# kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
myapp-5f9mv 1/1     Running   0          30s
myapp-bzxfh 1/1     Running   0          30s
myapp-czfbdb 1/1    Running   0          30s
myapp-kgwxv  1/1     Running   0          30s
myapp-kxxcw  1/1     Running   0          30s
myapp-m697l  1/1     Running   0          30s
myapp-p475k  1/1     Running   0          30s
myapp-r8pqp  1/1     Running   0          30s
myapp-s6htg  1/1     Running   0          30s
myapp-tv885  1/1     Running   0          30s
mypad       1/1     Running   0          19m
root@master-node:~# |
```

Yes! Since you're using a **ReplicationController**, it ensures that the desired number of replicas (Pods) is always running. If you manually delete a Pod, the ReplicationController will **automatically create a new Pod** to maintain the specified count.

```
root@master-node:~# kubectl delete pod myapp-5f9mv
pod "myapp-5f9mv" deleted
root@master-node:~#
root@master-node:~# kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
myapp-bzxfh 1/1     Running   0          6m11s
myapp-czfbdb 1/1    Running   0          6m11s
myapp-kgwxv  1/1     Running   0          6m11s
myapp-kxxcw  1/1     Running   0          6m11s
myapp-m697l  1/1     Running   0          6m11s
myapp-n2nps  1/1     Running   0          7s
myapp-p475k  1/1     Running   0          6m11s
myapp-r8pqp  1/1     Running   0          6m11s
myapp-s6htg  1/1     Running   0          6m11s
myapp-tv885  1/1     Running   0          6m11s
mypad       1/1     Running   0          24m
root@master-node:~# |
```

The **ReplicationController** continuously monitors the state of the cluster. If the number of running Pods is **less than the desired replicas (10 in your case)**, it will automatically create a new Pod to match the desired state.

## Delete RC

```
kubectl delete rc myapp
```

```
root@master-node:~# kubectl delete rc myapp
replicationcontroller "myapp" deleted
root@master-node:~#
root@master-node:~# kubectl get pods
NAME      READY     STATUS    RESTARTS   AGE
mypod    1/1      Running   0          40m
root@master-node:~# |
```

## 3. ReplicaSet – Let's perform some tasks using ReplicaSet

### Difference Between ReplicationController and ReplicaSet in Kubernetes

Both **ReplicationController** and **ReplicaSet** are used to ensure that a specified number of Pod replicas are running in Kubernetes. However, **ReplicaSet is the modern replacement for ReplicationController** with more advanced features.

## 1. ReplicationController (Old)

- Ensures a fixed number of Pod replicas
- Restarts missing Pods automatically
- Uses exact match selectors (less flexible)
- Deprecated in favor of ReplicaSet

### Example: ReplicationController YAML

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: myapp
spec:
```

```
replicas: 3
selector:
  app: webserver
template:
  metadata:
    labels:
      app: webserver
spec:
  containers:
    - name: nginx
      image: nginx
```

## 2. ReplicaSet (New & Improved)

- Improved version of ReplicationController**
- Supports "Set-Based" selectors** (more flexible)
- Works with Deployments** for better rollouts
- Still not commonly used directly (Deployments are preferred)**

### Example: ReplicaSet YAML

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-rs
spec:
  replicas: 3
  selector:
    matchLabels:
      app: webserver
  template:
    metadata:
      labels:
        app: webserver
    spec:
      containers:
```

```
- name: nginx
  image: nginx
```

## Let's create a ReplicaSet

```
vim rs.yaml
```

```
apiVersion: apps/v1
```

```
kind: ReplicaSet
```

```
metadata:
```

```
  name: mywebsite
```

```
spec:
```

```
  selector:
```

```
    matchLabels:
```

```
      app: webserver
```

```
  replicas: 20
```

```
  template:
```

```
    metadata:
```

```
      name: xyz
```

```
      labels:
```

```
        app: webserver
```

```
  spec:
```

```
    containers:
```

```
      - name: sdfdf
```

```
        image: httpd
```

```
root@master-node:~# vim rs.yaml
```

```
root@master-node:~# kubectl create -f rs.yaml
```

```
replicaset.apps/mywebsite created
```

```
root@master-node:~# kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
------	---------	---------	-------	-----

mywebsite	20	20	20	6s
-----------	----	----	----	----

```
root@master-node:~# kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
------	-------	--------	----------	-----

mypod	1/1	Running	0	3h4m
-------	-----	---------	---	------

mywebsite-24w5c	1/1	Running	0	17s
-----------------	-----	---------	---	-----

mywebsite-2hhv5	1/1	Running	0	17s
-----------------	-----	---------	---	-----

Okay, now let's change the desired state of the ReplicaSet. Here, we can **scale up and scale down** as needed. For example, if I currently have **20 pods** and I set the replicas to **15**, it will delete **5 pods**, leaving a total of **15 pods** running.

```
kubectl scale rs --replicas 15 mywebsite
```

```
root@master-node:~# kubectl scale rs --replicas 15 mywebsite
```

```
replicaset.apps/mywebsite scaled
```

```
root@master-node:~#
```

```
root@master-node:~# kubectl get rs
```

NAME	DESIRED	CURRENT	READY	AGE
------	---------	---------	-------	-----

mywebsite	15	15	15	4m57s
-----------	----	----	----	-------

```
root@master-node:~# |
```

We can scale up and scale down by modifying the file as well. **Let's do it.**

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: mywebsite
spec:
  selector:
    matchLabels:
      app: webserver
  replicas: 10
  template:
    metadata:
      name: xyz
      labels:
        app: webserver
    spec:
      containers:
        - name: sdfdf
          image: httpd
```

```
kubectl replace -f rs.yaml
```

```
root@master-node:~# kubectl replace -f rs.yaml
replicaset.apps/mywebsite replaced
root@master-node:~# kubectl get rs
NAME      DESIRED   CURRENT   READY   AGE
mywebsite  10        10        10      8m32s
root@master-node:~# ~^C
root@master-node:~# |
```

### 3. Deployment

Let's create a deployment.

```
vim deployment.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mywebsite
spec:
  selector:
    matchLabels:
      app: webserver
  replicas: 50
  template:
    metadata:
      name: xyz
      labels:
        app: webserver
    spec:
      containers:
        - name: sdfdf
          image: httpd
```

```
kubectl create -f deployment.yaml
```

```

root@master-node:~# kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite   50/50   50           50          2m45s
root@master-node:~#
root@master-node:~# kubectl get rs
NAME            DESIRED   CURRENT   READY   AGE
mywebsite-84ffd6ccdf   50        50        50      2m51s
root@master-node:~#

```

Here we can see here Deployment automatically create a replicaset.

Okay, now let's perform a **rollout**.

We will update the **nginx image version** in our deployment.yaml file, meaning we will deploy a different version.

In this process:

- The **old pods will terminate**.
- **New pods** with the updated image version will be created to match the desired state.

```

apiVersion: apps/v1
kind: Deployment
metadata:
  name: mywebsite
spec:
  selector:
    matchLabels:
      app: webserver
  replicas: 50
  template:
    metadata:
      name: xyz
      labels:
        app: webserver
    spec:
      containers:
        - name: sdfdf
          image: nginx:1.22

```

```
kubectl replace -f deployment.yaml
```

```
root@master-node:~# kubectl replace -f deployment.yaml
deployment.apps/mywebsite replaced
root@master-node:~# kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite   38/50   25          38          12m
root@master-node:~# kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite   38/50   50          38          12m
root@master-node:~# kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite   40/50   50          40          12m
root@master-node:~# |
```

No again rollout with nginx latest version

vim deployment.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: mywebsite
spec:
  selector:
    matchLabels:
      app: webserver
  replicas: 50
  template:
    metadata:
      name: xyz
      labels:
        app: webserver
    spec:
      containers:
        - name: sdfdf
          image: nginx:latest
```

kubectl replace -f deployment.yaml

```
root@master-node:~# kubectl replace -f deployment.yaml
deployment.apps/mywebsite replaced
root@master-node:~#
root@master-node:~# kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite   38/50   25          38          15m
root@master-node:~#
root@master-node:~# kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite   38/50   41          38          15m
root@master-node:~# kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite   38/50   50          38          15m
root@master-node:~# kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite   50/50   50          50          15m
root@master-node:~# vim deployment.yaml
root@master-node:~# |
```

```
root@master-node:~# kubectl rollout status deployment mywebsite
deployment "mywebsite" successfully rolled out
root@master-node:~#
root@master-node:~# |
```

### kubectl describe deployments.apps

```
root@master-node:~# kubectl describe deployments.apps
Name:           mywebsite
Namespace:      default
CreationTimestamp: Thu, 20 Feb 2025 17:39:48 +0000
Labels:          <none>
Annotations:    deployment.kubernetes.io/revision: 3
Selector:        app=webserver
Replicas:       50 desired | 50 updated | 50 total | 50 available | 0 unavailable
StrategyType:   RollingUpdate
MinReadySeconds: 0
RollingUpdateStrategy: 25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=webserver
  Containers:
    sdfdf:
      Image:      nginx:latest
      Port:       <none>
      Host Port: <none>
      Environment: <none>
```

### kubectl get rs

```
root@master-node:~# kubectl get rs
NAME           DESIRED   CURRENT   READY   AGE
mywebsite-55d8c74676   0         0         0      7m56s
mywebsite-5979986654   50        50        50     5m1s
mywebsite-84ffd6ccdf   0         0         0      20m
root@master-node:~# |
```

kubectl rollout history deployment mywebsite

```
root@master-node:~# kubectl rollout history deployment mywebsite
deployment.apps/mywebsite
REVISION  CHANGE-CAUSE
1          <none>
2          <none>
3          <none>
```

```
root@master-node:~# |
```

kubectl set image deployment mywebsite nginx=nginx:latest --record

```
root@master-node:~# kubectl set image deployment mywebsite nginx=nginx:latest --record
Flag --record has been deprecated, --record will be removed in the future
deployment.apps/mywebsite image updated
error: unable to find container named "nginx"
root@master-node:~# kubectl rollout history deployment mywebsite
deployment.apps/mywebsite
REVISION  CHANGE-CAUSE
1          <none>
2          <none>
3          kubectl set image deployment mywebsite nginx=nginx:latest --record=true
```

```
root@master-node:~# |
```

kubectl rollout history deployment mywebsite --revision 1

```
root@master-node:~# kubectl rollout history deployment mywebsite --revision 1
deployment.apps/mywebsite with revision #1
Pod Template:
  Labels:      app=webserver
                pod-template-hash=84ffd6ccdf
  Containers:
    sdfdf:
      Image:      httpd
      Port:       <none>
      Host Port: <none>
      Environment: <none>
      Mounts:     <none>
      Volumes:   <none>
      Node-Selectors: <none>
      Tolerations:  <none>
root@master-node:~# |
```

[kubectl rollout undo deployment mywebsite --to-revision 1](#)

```
root@master-node:~# kubectl rollout undo deployment mywebsite --to-revision 1
deployment.apps/mywebsite rolled back
root@master-node:~#
root@master-node:~# kubectl get deployments
NAME        READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite   38/50   48           38          31m
root@master-node:~# kubectl get pods
NAME                           READY   STATUS    RESTARTS   AGE
mywebsite-84ffd6ccdf-2wfqp   1/1     Running   0          14s
mywebsite-84ffd6ccdf-4gcbd   1/1     Running   0          14s
mywebsite-84ffd6ccdf-4l7g4   1/1     Running   0          13s
```

[kubectl rollout history deployment mywebsite](#)

```
root@master-node:~# kubectl rollout history deployment mywebsite
deployment.apps/mywebsite
REVISION  CHANGE-CAUSE
2        <none>
3        kubectl set image deployment mywebsite nginx=nginx:latest --record=true
4        <none>
root@master-node:~# |
```

[kubectl scale deployment mywebsite --replicas 60](#)

```
root@master-node:~# kubectl scale deployment mywebsite --replicas 60
deployment.apps/mywebsite scaled
root@master-node:~#
root@master-node:~# kubectl get deployments
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite  60/60   60          60          36m
root@master-node:~# |
```

kubectl edit deployments.apps mywebsite

```
metadata:
  annotations:
    deployment.kubernetes.io/revision: "4"
  creationTimestamp: "2025-02-20T17:39:48Z"
  generation: 6
  name: mywebsite
  namespace: default
  resourceVersion: "75167"
  uid: c1aa0879-3698-4aa4-aaf1-4cace888ac5c
spec:
  progressDeadlineSeconds: 600
  replicas: 75
  revisionHistoryLimit: 10
  selector:
    matchLabels:
      app: webserver
root@master-node:~# kubectl edit deployments.apps mywebsite
deployment.apps/mywebsite edited
root@master-node:~#
root@master-node:~# kubectl get deployments
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite  64/75   75          64          41m
root@master-node:~# kubectl get deployments
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
mywebsite  68/75   75          68          41m
```

```
root@master-node:~# kubectl delete deployments.apps mywebsite
deployment.apps "mywebsite" deleted
root@master-node:~#
root@master-node:~# kubectl get deployments.apps
No resources found in default namespace.
root@master-node:~#
root@master-node:~# kubectl get rs
No resources found in default namespace.
root@master-node:~# |
```

kubectl create deployment --image httpd --replicas 7 test

```
root@master-node:~# kubectl create deployment --image httpd --replicas 7 test
deployment.apps/test created
root@master-node:~#
root@master-node:~# kubectl get deployments
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
test      7/7     7           7           6s
root@master-node:~# kubectl get rs
NAME          DESIRED   CURRENT   READY   AGE
test-6866fd55cd   7         7         7         9s
root@master-node:~# |
```

[kubectl create deployment web-server --replicas 20 --image nginx --dry-run=client -o yaml](#)

```
root@master-node:~# kubectl create deployment web-server --replicas 20 --image nginx --dry-run=client -o yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: web-server
  name: web-server
spec:
  replicas: 20
  selector:
    matchLabels:
      app: web-server
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: web-server
    spec:
      containers:
      - image: nginx
        name: nginx
        resources: {}
status: {}
root@master-node:~# |
```

[kubectl create deployment web-server --replicas 20 --image nginx --dry-run=client -o yaml > nginx.yaml](#)

[kubectl run neel --image httpd --dry-run=client -o yaml](#)

```
root@master-node:~# kubectl run neel --image httpd --dry-run=client -o yaml
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: neel
    name: neel
spec:
  containers:
  - image: httpd
    name: neel
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
root@master-node:~# |
```

kubectl run neel --image httpd --dry-run=client -o yaml > web.yaml